

Getting started with the Grid Application Toolkit

Rob van Nieuwpoort

rob@cs.vu.nl

Vrije Universiteit
Amsterdam



- What is GAT and why do we need it?
- JavaGAT structure and overview
- Security
- Grid I/O
- Resource Management
- Application Information Management
- Monitoring
- Hands-on session

- EU Project Funded by 5th Framework
- 3 years, 5 million euros
- Partners
 - PSNC, AEI, ZIB, MASARYK, SZTAKI, ISUFI, Cardiff, NTUA, Chicago, ISI, Wisconsin, Sun, Compaq,...
 - Vrije Universiteit Amsterdam
- 12 Work Packages covering
 - Grid Portals, Mobile Users, Data Management, Applications, Testbed, ...
 - GAT: The Grid Application Toolkit

Grid Application

`File.copy(...)`



Grid Application


`submitJob(...)`



`File.copy(...)`

Grid Application


`submitJob(...)`



`File.copy(...)`

Grid Application

`submitJob(...)`




`cp`
`ftp`

`File.copy(...)`

Grid Application

`submitJob(...)`




cp
ftp
gridftp

`File.copy(...)`

Grid Application

`submitJob(...)`




`cp`
`ftp`
`gridftp`
`scp`

`File.copy(...)`

Grid Application

`submitJob(...)`




`cp`
`ftp`
`gridftp`
`scp`
`http`




`File.copy(...)`

Grid Application


`submitJob(...)`




- Which should you use?




- Which should you use?
- Some might not be available on all sites




- Which should you use?
- Some might not be available on all sites
- Some may not work for all users (certificates)



- Which should you use?
- Some might not be available on all sites
- Some may not work for all users (certificates)
- Version differences (Globus changes every 2 months)



- Which should you use?
- Some might not be available on all sites
- Some may not work for all users (certificates)
- Version differences (Globus changes every 2 months)
- Exponential number of combinations!



- Which should you use?
- Some might not be available on all sites
- Some may not work for your testing context
- Versions differ a lot (Globus changes every 2 months)
- Exponential number of combinations!

Not portable!

```
int RemoteFile::GetFile (char const* source,  char const* target) {
    globus_url_t          source_url;
    globus_io_handle_t     dest_io_handle;
    globus_ftp_client_operationattr_t  source_ftp_attr;
    globus_result_t        result;
    globus_gass_transfer_requestattr_t source_gass_attr;
    globus_gass_copy_attr_t source_gass_copy_attr;
    globus_gass_copy_handle_t gass_copy_handle;
    globus_gass_copy_handleattr_t gass_copy_handleattr;
    globus_ftp_client_handleattr_t ftp_handleattr;
    globus_io_attr_t        io_attr;
    int                     output_file = -1;

    if ( globus_url_parse (source_URL, &source_url) != GLOBUS_SUCCESS ) {
        printf ("can not parse source_URL \"%s\"\n", source_URL);
        return (-1);
    }

    if ( source_url.scheme_type != GLOBUS_URL_SCHEME_GSIFTP &&
        source_url.scheme_type != GLOBUS_URL_SCHEME_FTP &&
        source_url.scheme_type != GLOBUS_URL_SCHEME_HTTP &&
        source_url.scheme_type != GLOBUS_URL_SCHEME_HTTPS ) {
        printf ("can not copy from %s - wrong prot\n", source_URL);
        return (-1);
    }

    globus_gass_copy_handleattr_init (&gass_copy_handleattr);
    globus_gass_copy_attr_init (&source_gass_copy_attr);

    globus_ftp_client_handleattr_init (&ftp_handleattr);
    globus_io_fileattr_init (&io_attr);

    globus_gass_copy_attr_set_io (&source_gass_copy_attr, &io_attr);
    globus_gass_copy_handleattr_set_ftp_attr (&gass_copy_handleattr,
        &ftp_handleattr);

    globus_gass_copy_handle_init (&gass_copy_handle,
        &gass_copy_handleattr);
}
```

```
if (source_url.scheme_type == GLOBUS_URL_SCHEME_GSIFTP ||
    source_url.scheme_type == GLOBUS_URL_SCHEME_FTP ) {
    globus_ftp_client_operationattr_init (&source_ftp_attr);
    globus_gass_copy_attr_set_ftp (&source_gass_copy_attr,
        &source_ftp_attr);
}
else {
    globus_gass_transfer_requestattr_init (&source_gass_attr,
        source_url.scheme);
    globus_gass_copy_attr_set_gass (&source_gass_copy_attr,
        &source_gass_attr);
}

output_file = globus_libc_open ((char*) target,
    O_WRONLY | O_TRUNC | O_CREAT,
    S_IRUSR | S_IWUSR | S_IRGRP |
    S_IWGRP);

if ( output_file == -1 ) {
    printf ("could not open the file \"%s\"\n", target);
    return (-1);
}

/* convert stdout to be a globus_io_handle */
if ( globus_io_file_posix_convert (output_file, 0,
        &dest_io_handle)
    != GLOBUS_SUCCESS) {
    printf ("Error converting the file handle\n");
    return (-1);
}

result = globus_gass_copy_register_url_to_handle (
    &gass_copy_handle, (char*)source_URL,
    &source_gass_copy_attr, &dest_io_handle,
    my_callback, NULL);

if ( result != GLOBUS_SUCCESS ) {
    printf ("error: %s\n", globus_object_printable_to_string
        (globus_error_get (result)));
    return (-1);
}

globus_url_destroy (&source_url);
return (0);
}
```

```
package org.globus.ogsa.gui;

import java.io.BufferedReader;
import java.io.File;
import java.io.FileReader;
import java.net.URL;
import java.util.Date;
import java.util.Vector;
import javax.xml.rpc.Stub;
import org.apache.axis.message.MessageElement;
import org.apache.axis.utils.XMLUtils;
import org.globus.*
import org.gridforum.ogsi.*
import org.gridforum.ogsi.holders.TerminationTimeTypeHolder;
import org.w3c.dom.Document;
import org.w3c.dom.Element;

public class RFTClient {
    public static void copy (String source_url, String target_url) {
        try {
            File requestFile = new File (source_url);
            BufferedReader reader = null;
            try {
                reader = new BufferedReader (new FileReader (requestFile));
            } catch (java.io.FileNotFoundException fnfe) { }
            Vector requestData = new Vector ();
            requestData.add (target_url);
            TransferType[] transfers1 = new TransferType[transferCount];
            RFTOptionsType multirftOptions = new RFTOptionsType ();

            multirftOptions.setBinary (Boolean.valueOf (
                (String)requestData.elementAt (0)).booleanValue ());
            multirftOptions.setBlockSize (Integer.valueOf (
                (String)requestData.elementAt (1)).intValue ());
            multirftOptions.setTcpBufferSize (Integer.valueOf (
                (String)requestData.elementAt (2)).intValue ());
            multirftOptions.setNotpt (Boolean.valueOf (
                (String)requestData.elementAt (3)).booleanValue ());
            multirftOptions.setParallelStreams (Integer.valueOf (
                (String)requestData.elementAt (4)).intValue ());
            multirftOptions.setDcau(Boolean.valueOf(
                (String)requestData.elementAt (5)).booleanValue ());

            int i = 7;
            for (int j = 0; j < transfers1.length; j++)
            {
                transfers1[j] = new TransferType ();

                transfers1[j].setTransferId (j);
                transfers1[j].setSourceUrl ((String)requestData.elementAt (i++));
                transfers1[j].setDestinationUrl ((String)requestData.elementAt (i++));
                transfers1[j].setRftOptions (multirftOptions);
            }
        }
    }
}
```

```
TransferRequestType transferRequest = new TransferRequestType ();
transferRequest.setTransferArray (transfers1);

int concurrency = Integer.valueOf
    ((String)requestData.elementAt(6)).intValue();

if (concurrency > transfers1.length)
{
    System.out.println ("Concurrency should be less than the number"
        "of transfers in the request");
    System.exit (0);
}
transferRequest.setConcurrency (concurrency);

TransferRequestElement requestElement = new TransferRequestElement ();
requestElement.setTransferRequest (transferRequest);

ExtensibilityType extension = new ExtensibilityType ();
extension = AnyHelper.getExtensibility (requestElement);

OGSIServiceGridLocator factoryService = new OGSIServiceGridLocator ();
Factory factory = factoryService.getFactoryPort (new URL (source_url));
GridServiceFactory gridFactory = new GridServiceFactory (factory);

LocatorType locator = gridFactory.createService (extension);
System.out.println ("Created an instance of Multi-RFT");

MultiFileRFTDefinitionServiceGridLocator loc
    = new MultiFileRFTDefinitionServiceGridLocator();
RFTPortType rftPort = loc.getMultiFileRFTDefinitionPort (locator);
((Stub)rftPort)._setProperty (Constants.AUTHORIZATION,
    NoAuthorization.getInstance());
((Stub)rftPort)._setProperty (GSIConstants.GSI_MODE,
    GSIConstants.GSI_MODE_FULL_DELEG);
((Stub)rftPort)._setProperty (Constants.GSI_SEC_CONV,
    Constants.SIGNATURE);
((Stub)rftPort)._setProperty (Constants.GRIM_POLICY_HANDLER,
    new IgnoreProxyPolicyHandler ());

int requestid = rftPort.start ();
System.out.println ("Request id: " + requestid);

}
catch (Exception e)
{
    System.err.println (MessageUtils.toString (e));
}
}
```

- The situation today:
 - Grids: everywhere
 - Grid applications: nowhere
- Why is this?
 - Application programmers accept the Grid as a computing paradigm only very slowly.
- Problems:
 - Interfaces are NOT simple
 - Portability / interoperability
 - Different and evolving interfaces to the 'Grid'
 - Environment changes in many ways
 - WSDL and web services do not solve all these problems

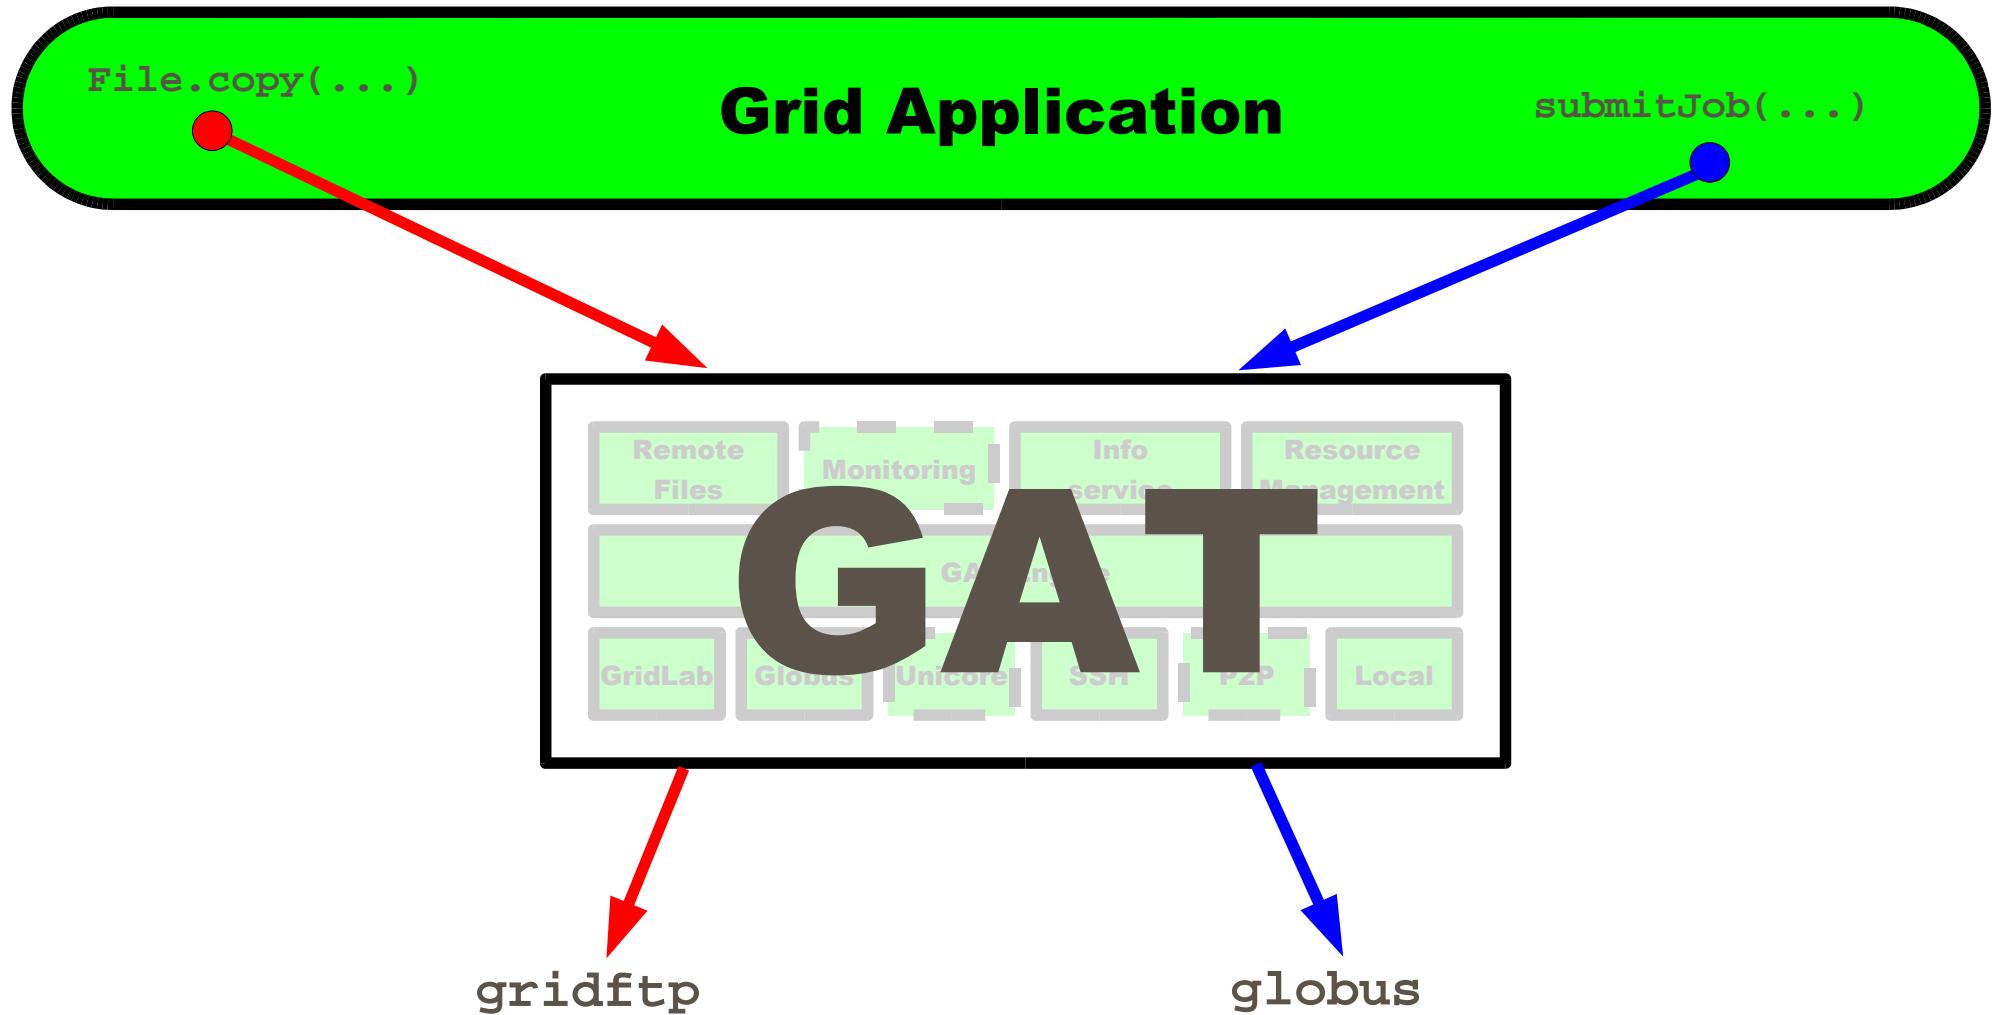
- GAT: Grid Application Toolkit
 - API and Toolkit for developing and running portable grid applications independently of the underlying grid infrastructure and available services
- GAT is used by applications to access grid services
- **Simple** API
- GAT Adaptors (“plugins”)
 - Connect GAT to grid services
 - Allow for multiple providers (Globus, Unicore, ProActive, ...)
- GAT Engine
 - Provides runtime delegation of GAT-API calls to adaptors

- GAT does not aim to replace existing “grid infrastructure.”
- GAT aims to provide a simple, clear interface to many different infrastructures
 - GRAM
 - Condor
 - Unicore
 - GridFTP
 - ...
- Open source, BSD-like licence

- Applications make GAT-API calls for grid operations
 - Applications link against GAT
- Applications run irrespective of available infrastructure
 - GAT Engine loads all available adaptors **at runtime**
 - Upon a call to the GAT-API the GAT Engine determines which adaptor(s) provide the “grid operation”
 - Upon “grid operation” failure another adaptor may be called
- There exist a set of default adaptors which provide default local capabilities
 - Grid applications can thus be compiled, linked, and tested without any available grid services
 - The same application executable can run in a “full grid environment.” **No recompilation / linking**

- A language independent, object oriented specification
- Language bindings
 - C, C++, Java, python, ...
- Work in progress
 - Perl, .net, fortran
- Adopt the “look-and-feel” of the target language
- focus on well-known programming paradigms
 - Provide a File API for file access
 - Not (web)services to services to files. . .
 - Programmers expect open, close, read, write, seek. Do not introduce fancy things like the need to ask a service discovery service to tell me the location of an service which is able to tell me the location of my file...

- Security (deal with passwords, credentials, etc)
- Grid I/O
 - File operations, remote file access, file replication
 - Inter-process communication
- Resource Management
 - Resource brokering
 - Forking grid applications, job management
- Application Information Management
 - Global repository for application specific information
 - Query this information repository
- Monitoring
 - Grid monitoring
 - Application monitoring and steering



File Copy with GAT (C++)

```
#include <GAT++.hpp>

void RemoteFile::GetFile (
    GAT::Context context, std::string source_url,
    std::string target_url) throws GAT::Exception {
    GAT::File file (context, source_url);
    file.Copy      (target_url);
}
```

File Copy with GAT (C++)

```
#include <GAT++.hpp>

void RemoteFile::GetFile (
    GAT::Context context, std::string source_url,
    std::string target_url) throws GAT::Exception {

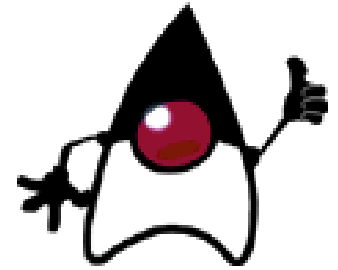
    GAT::File file (context, source_url);
    file.Copy      (target_url);
}
```

- What is GAT and why do we need it?
- **JavaGAT overview and structure**
- Security
- Grid I/O
- Resource Management
- Application Information Management
- Monitoring

- Hands-on session

Why Java?

- Java is widely used, object oriented.
- Secure (language, sandbox)
- Java is “write once, run everywhere”.
 - Compile application on your desktop machine.
 - This creates machine independent bytecode.
 - Copy application files and the GAT to any grid site (portal typically does this for you).
 - Just run it. No recompilation / configuration.
- Performance of current JITs is good.
 - Compiled (just-in-time)
 - Runtime, profile-driven optimizations
 - Applications are typically 10% slower than C.
- Ideal for grid computing?



- GAT API is an object oriented specification, but language independent.
- Tension between existing language features and the GAT specification.
 - Programmers do not want to learn new APIs for features already in the language.
 - GAT should “feel” the same across platforms.
- Find the tight trade-off

Grid Application

Files

Monitoring

**Info
service**

**Resource
Management**

GAT Engine

GridLab

Globus

Unicore

SSH

P2P

Local

Legend:

Java

Done

W.I.P

Grid Application

API

Files

Monitoring

Info
service

Resource
Management

GAT Engine

GridLab

Globus

Unicore

SSH

P2P

Local

Legend:

Java

Done

W.I.P

Grid Application

API

Files

Monitoring

Info
service

Resource
Management

GAT Engine

ADAPT.

GridLab

Globus

Unicore

SSH

P2P

Local

Legend:

Java

Done

W.I.P

`File.copy(...)`



Grid Application

API

Files

Monitoring

Info
service

Resource
Management

GAT Engine

ADAPT.

GridLab

Globus

Unicore

SSH

P2P

Local

Legend:

Java

Done

W.I.P

`File.copy(...)`

Grid Application

Files

Monitoring

Info
service

Resource
Management

GAT Engine

GridLab

Globus

Unicore

SSH

P2P

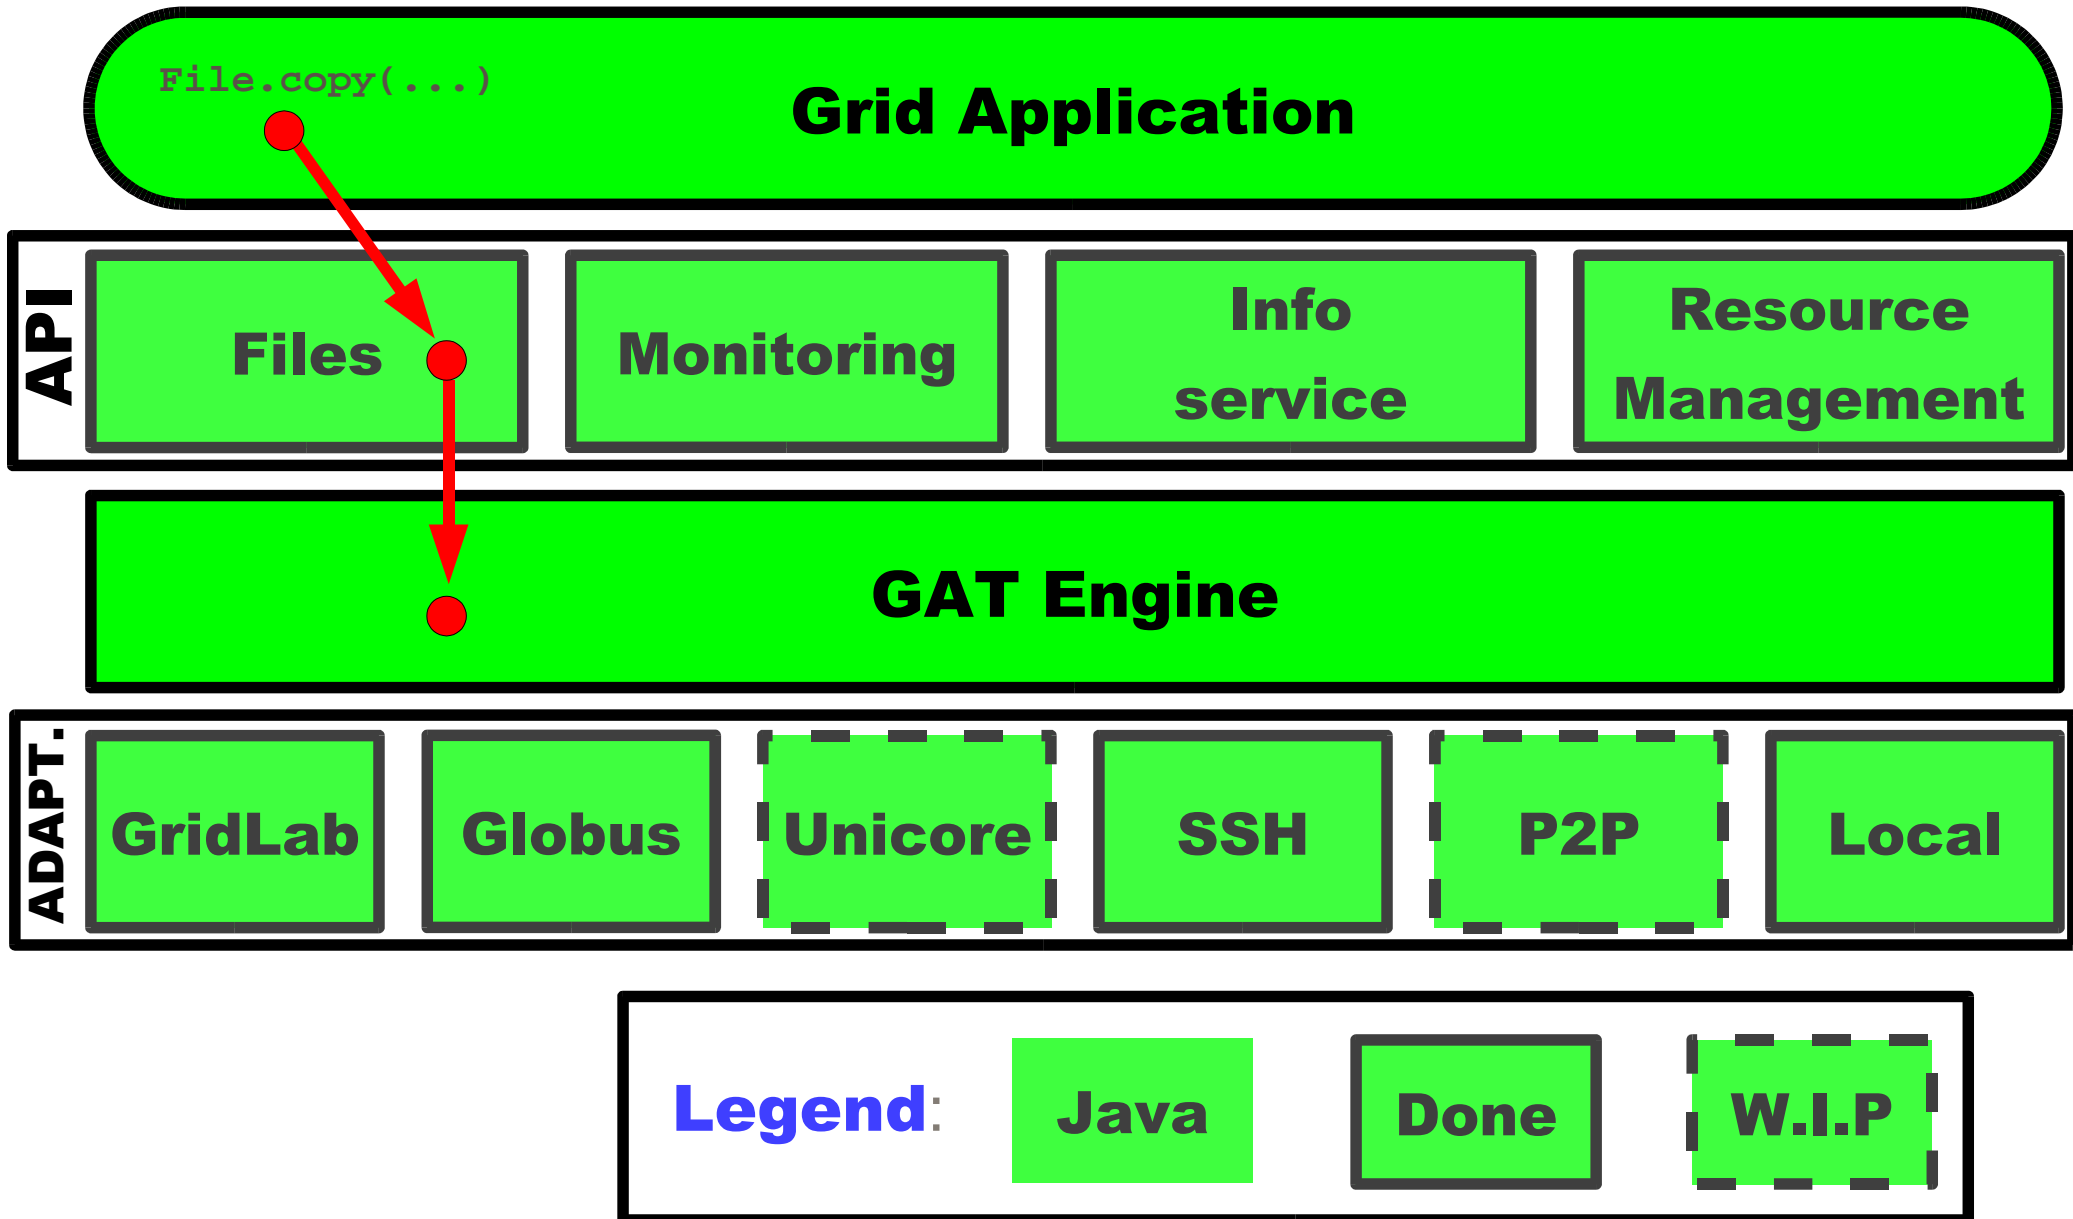
Local

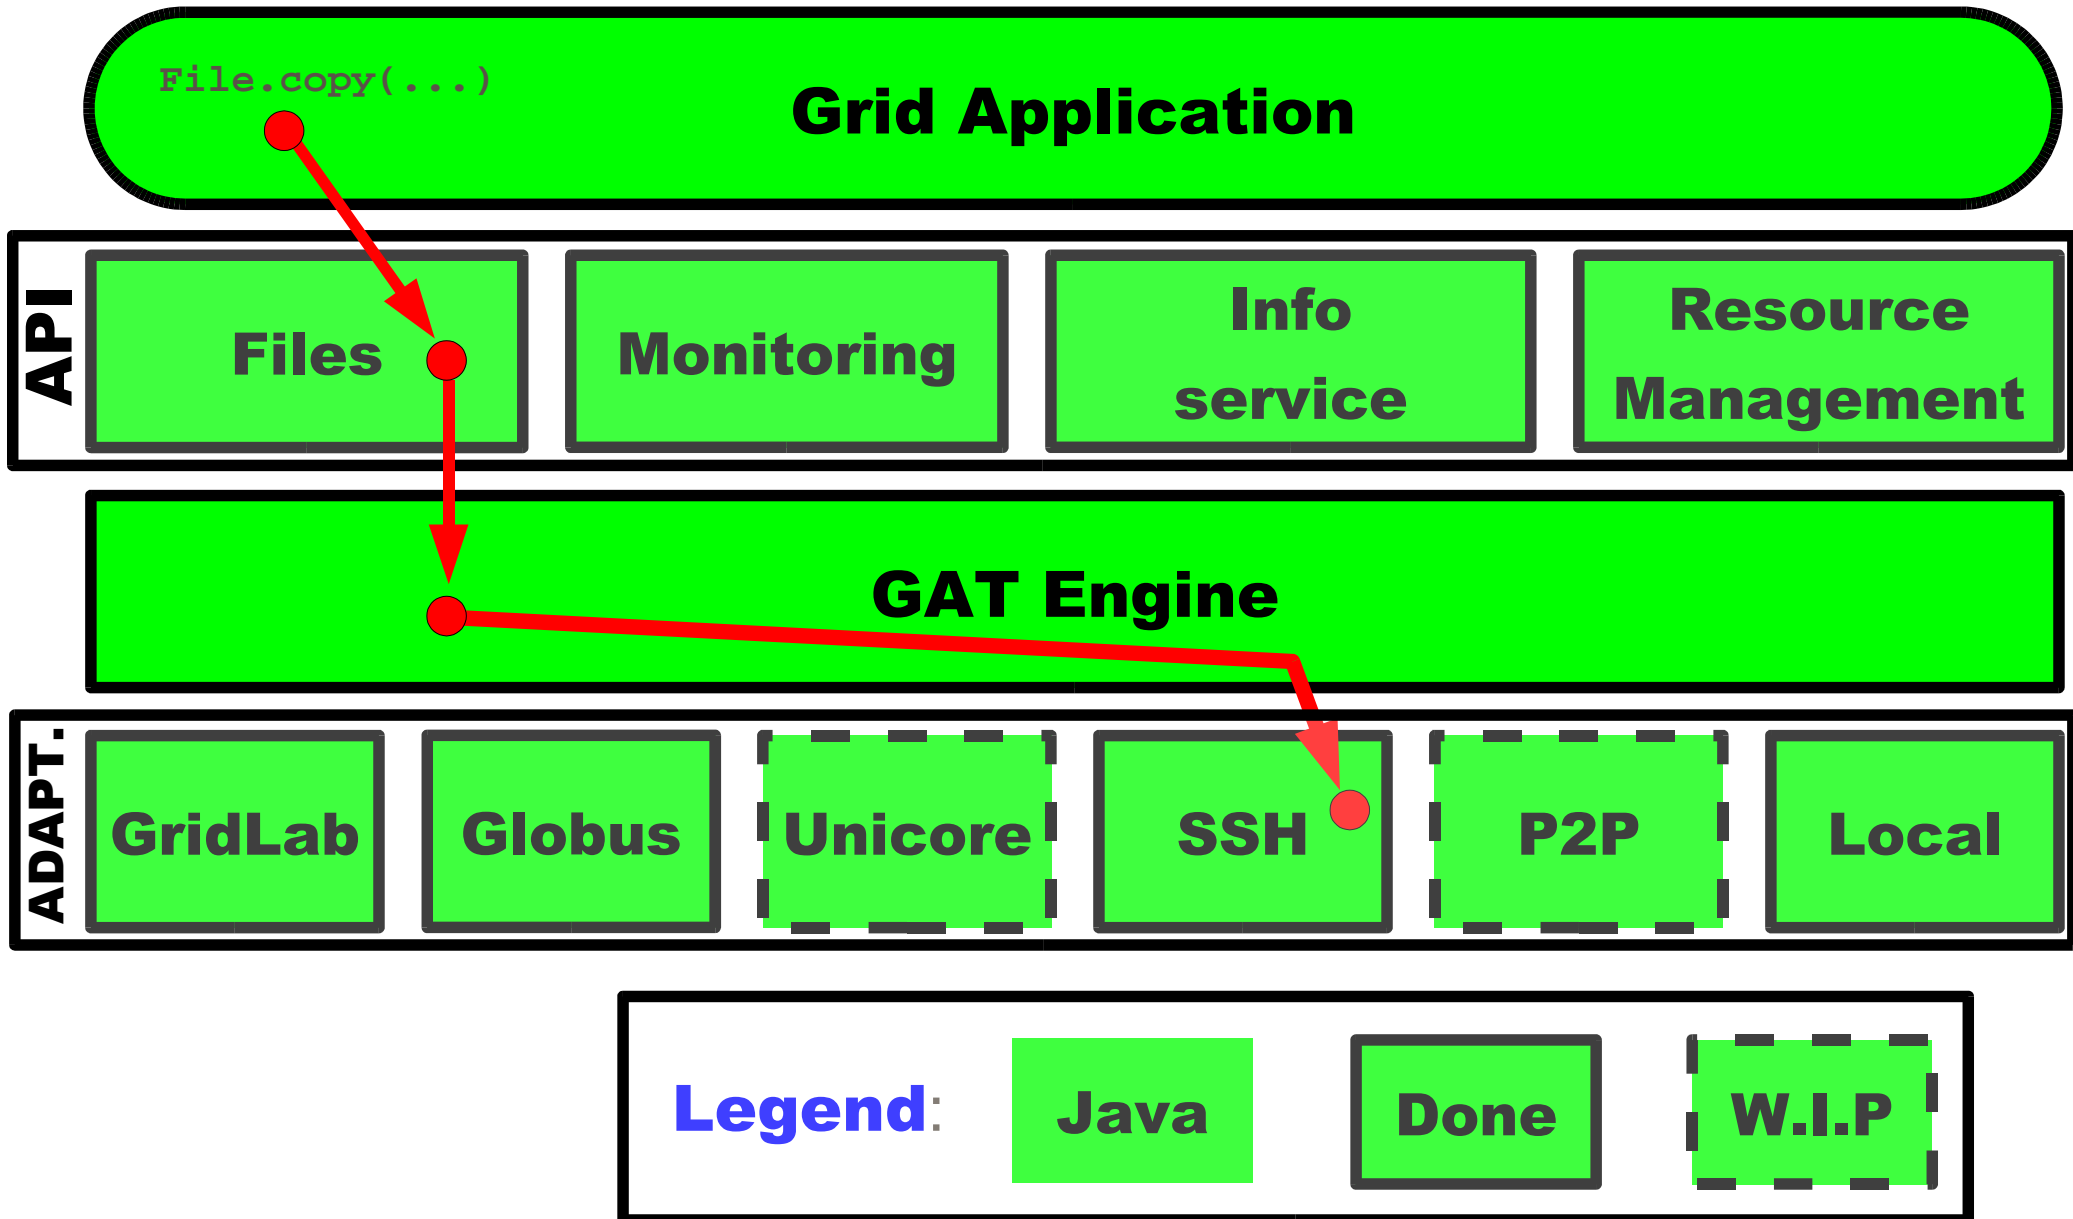
Legend:





Java

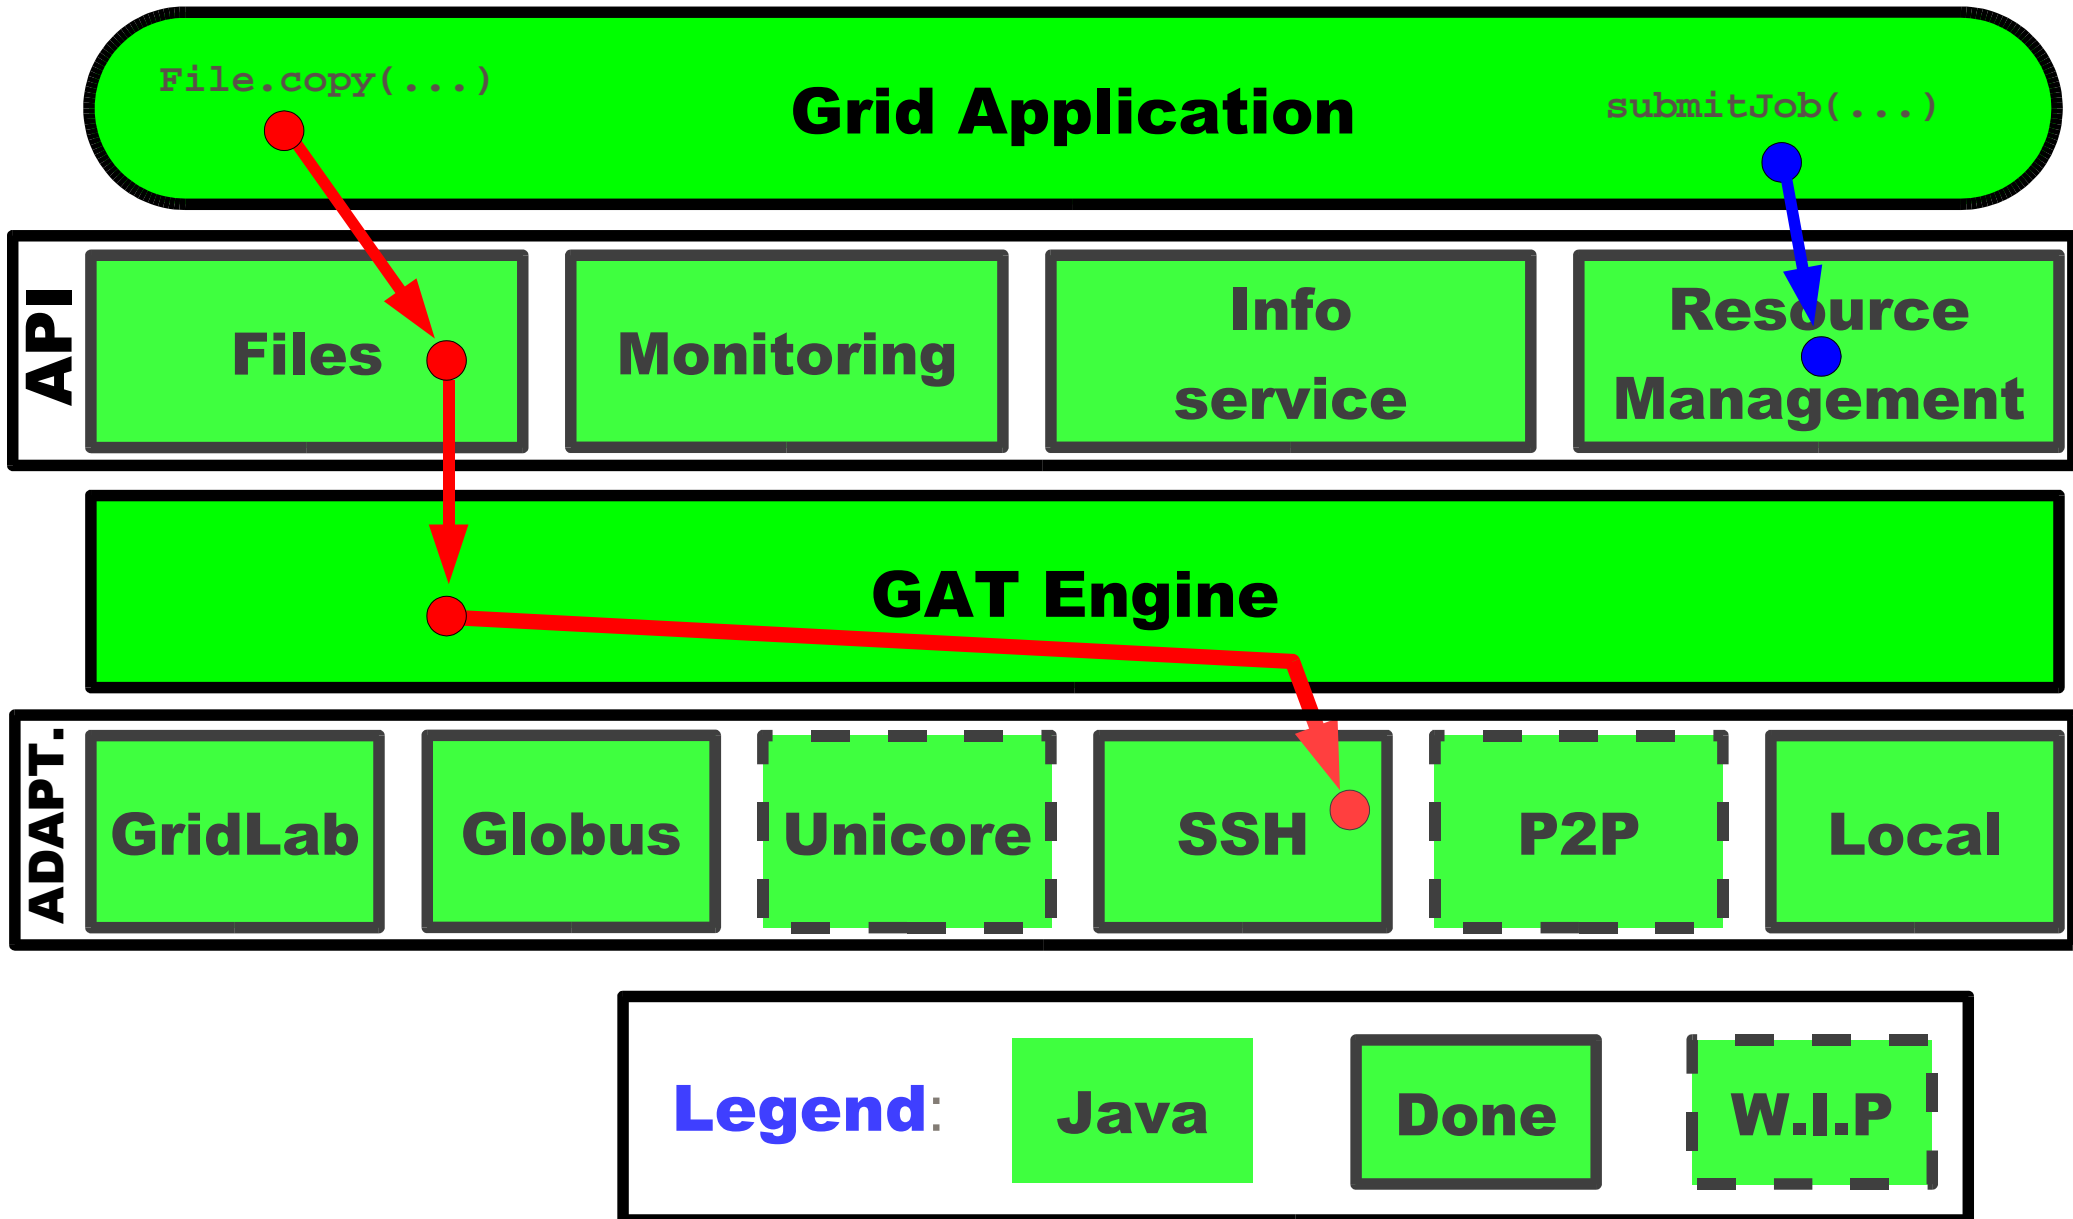
Done

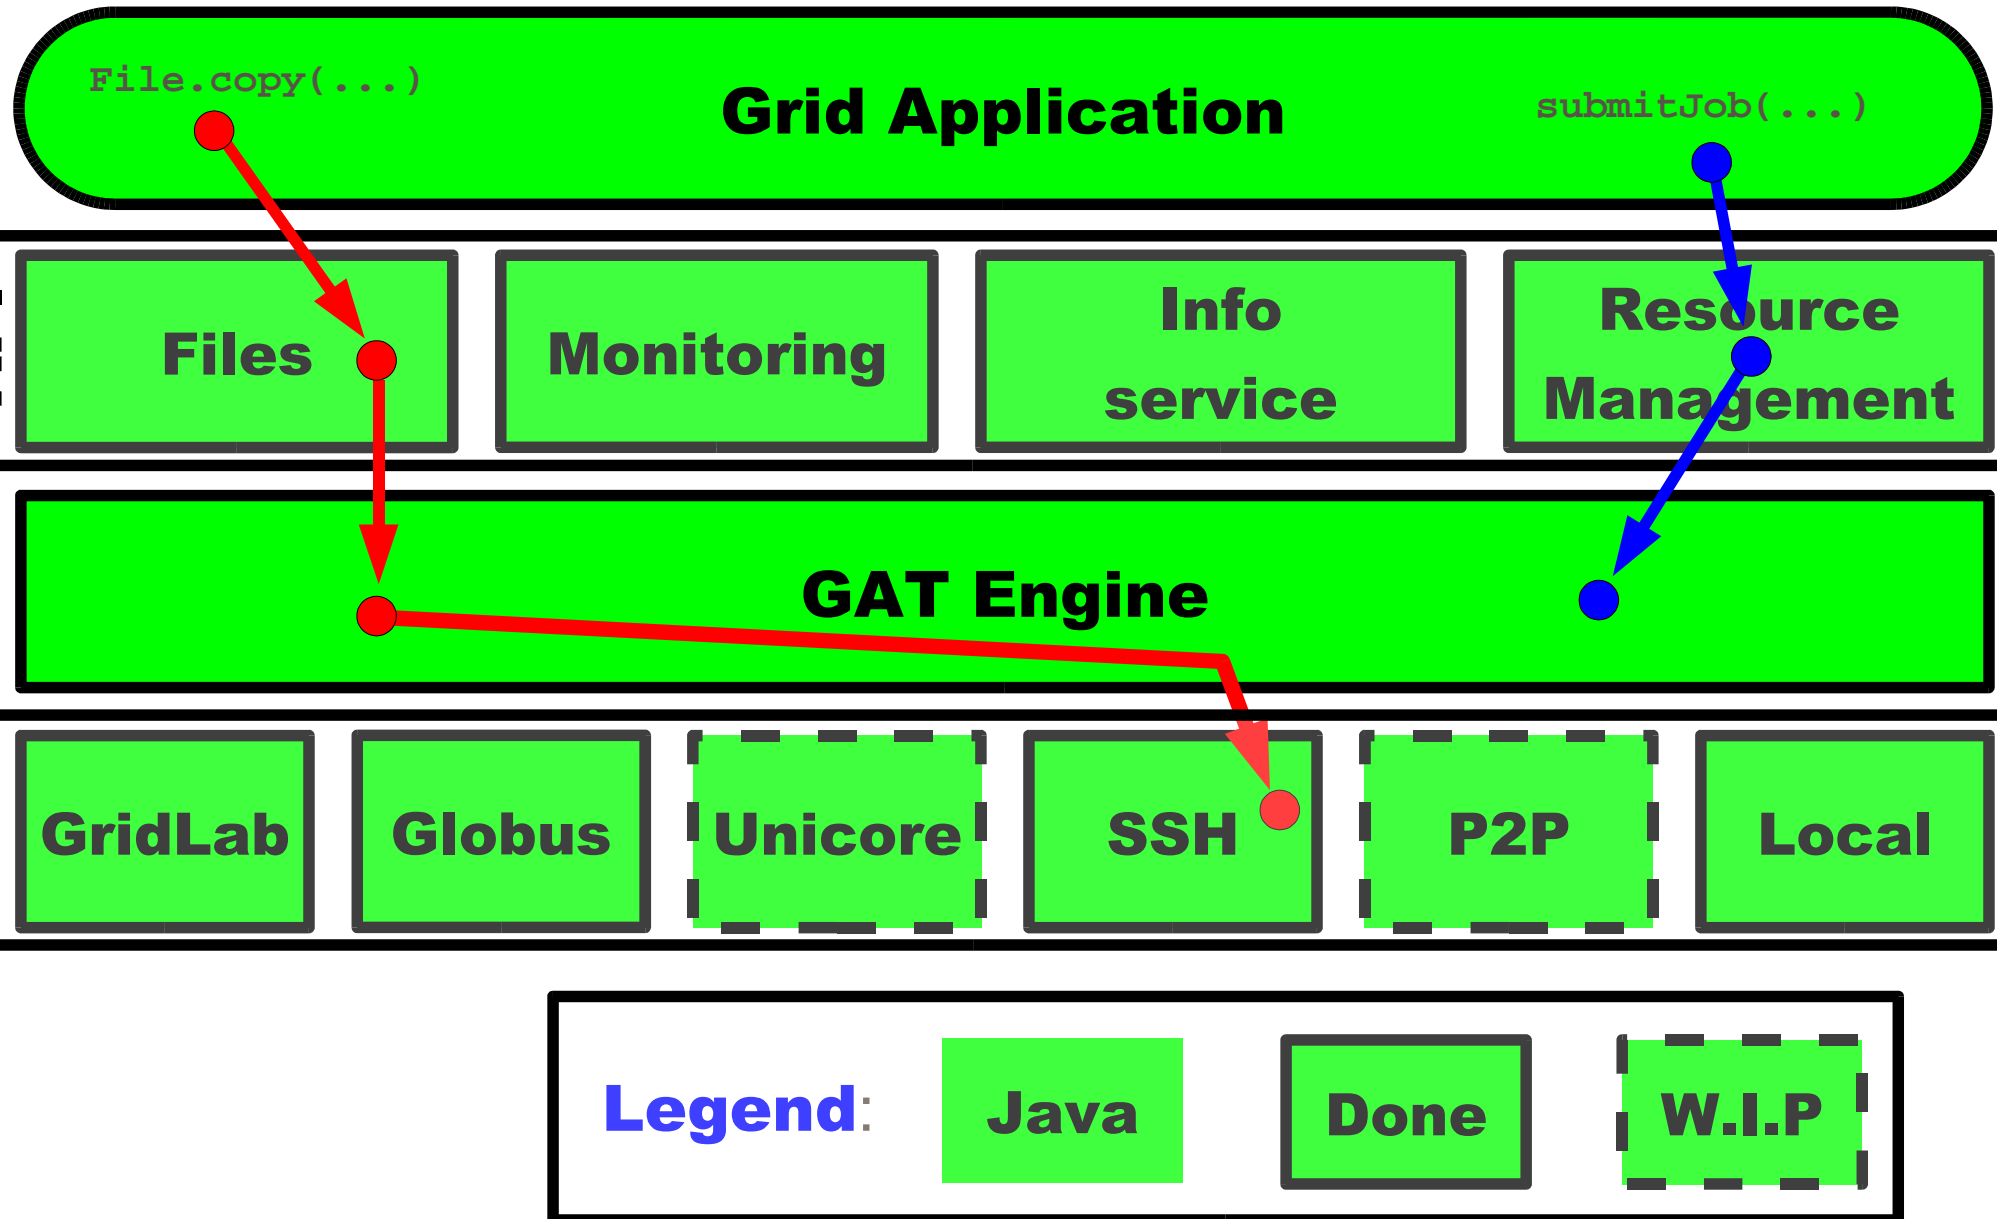
W.I.P

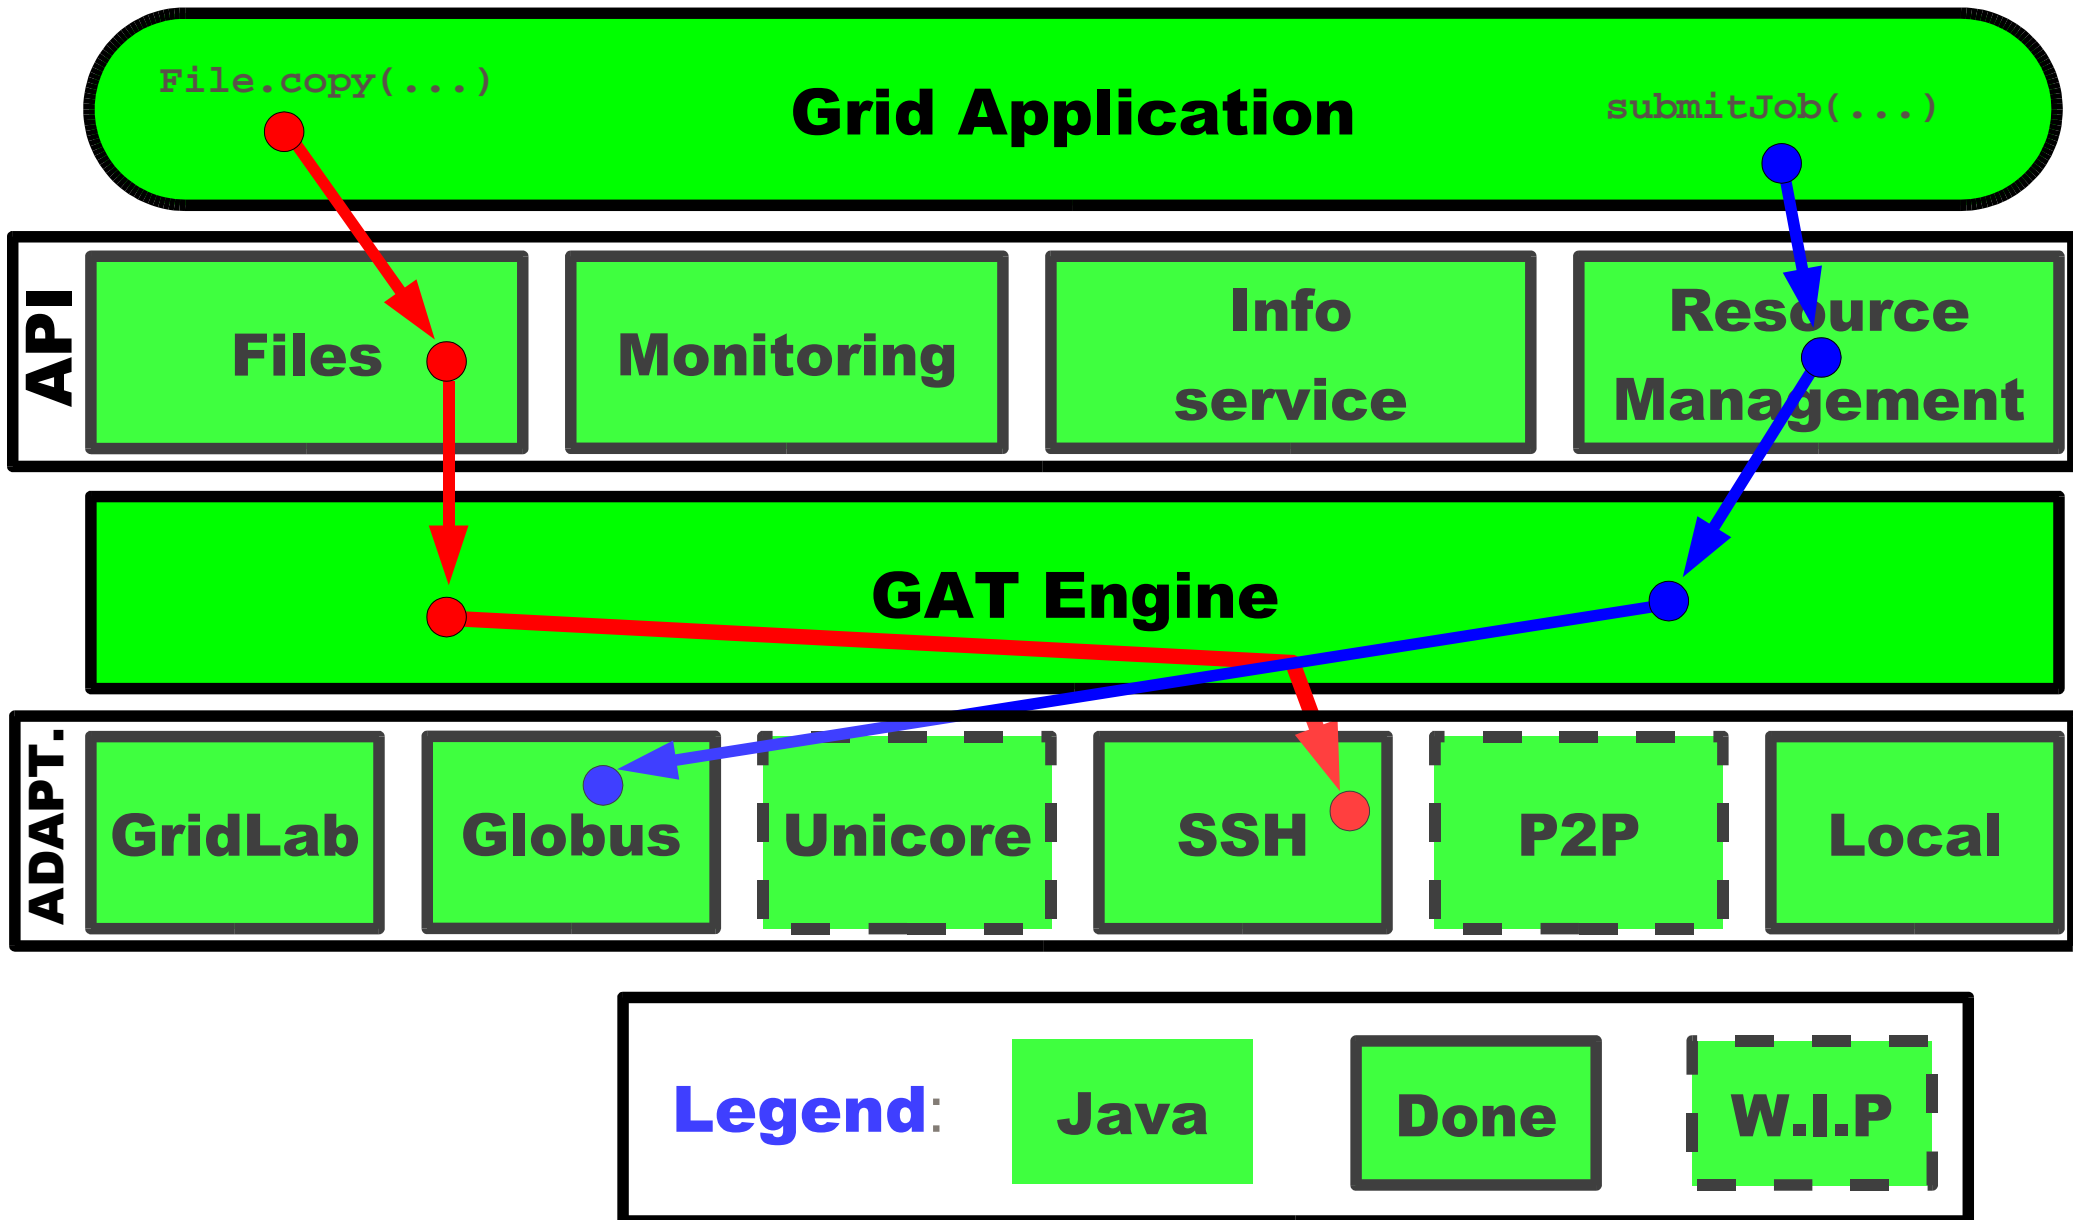


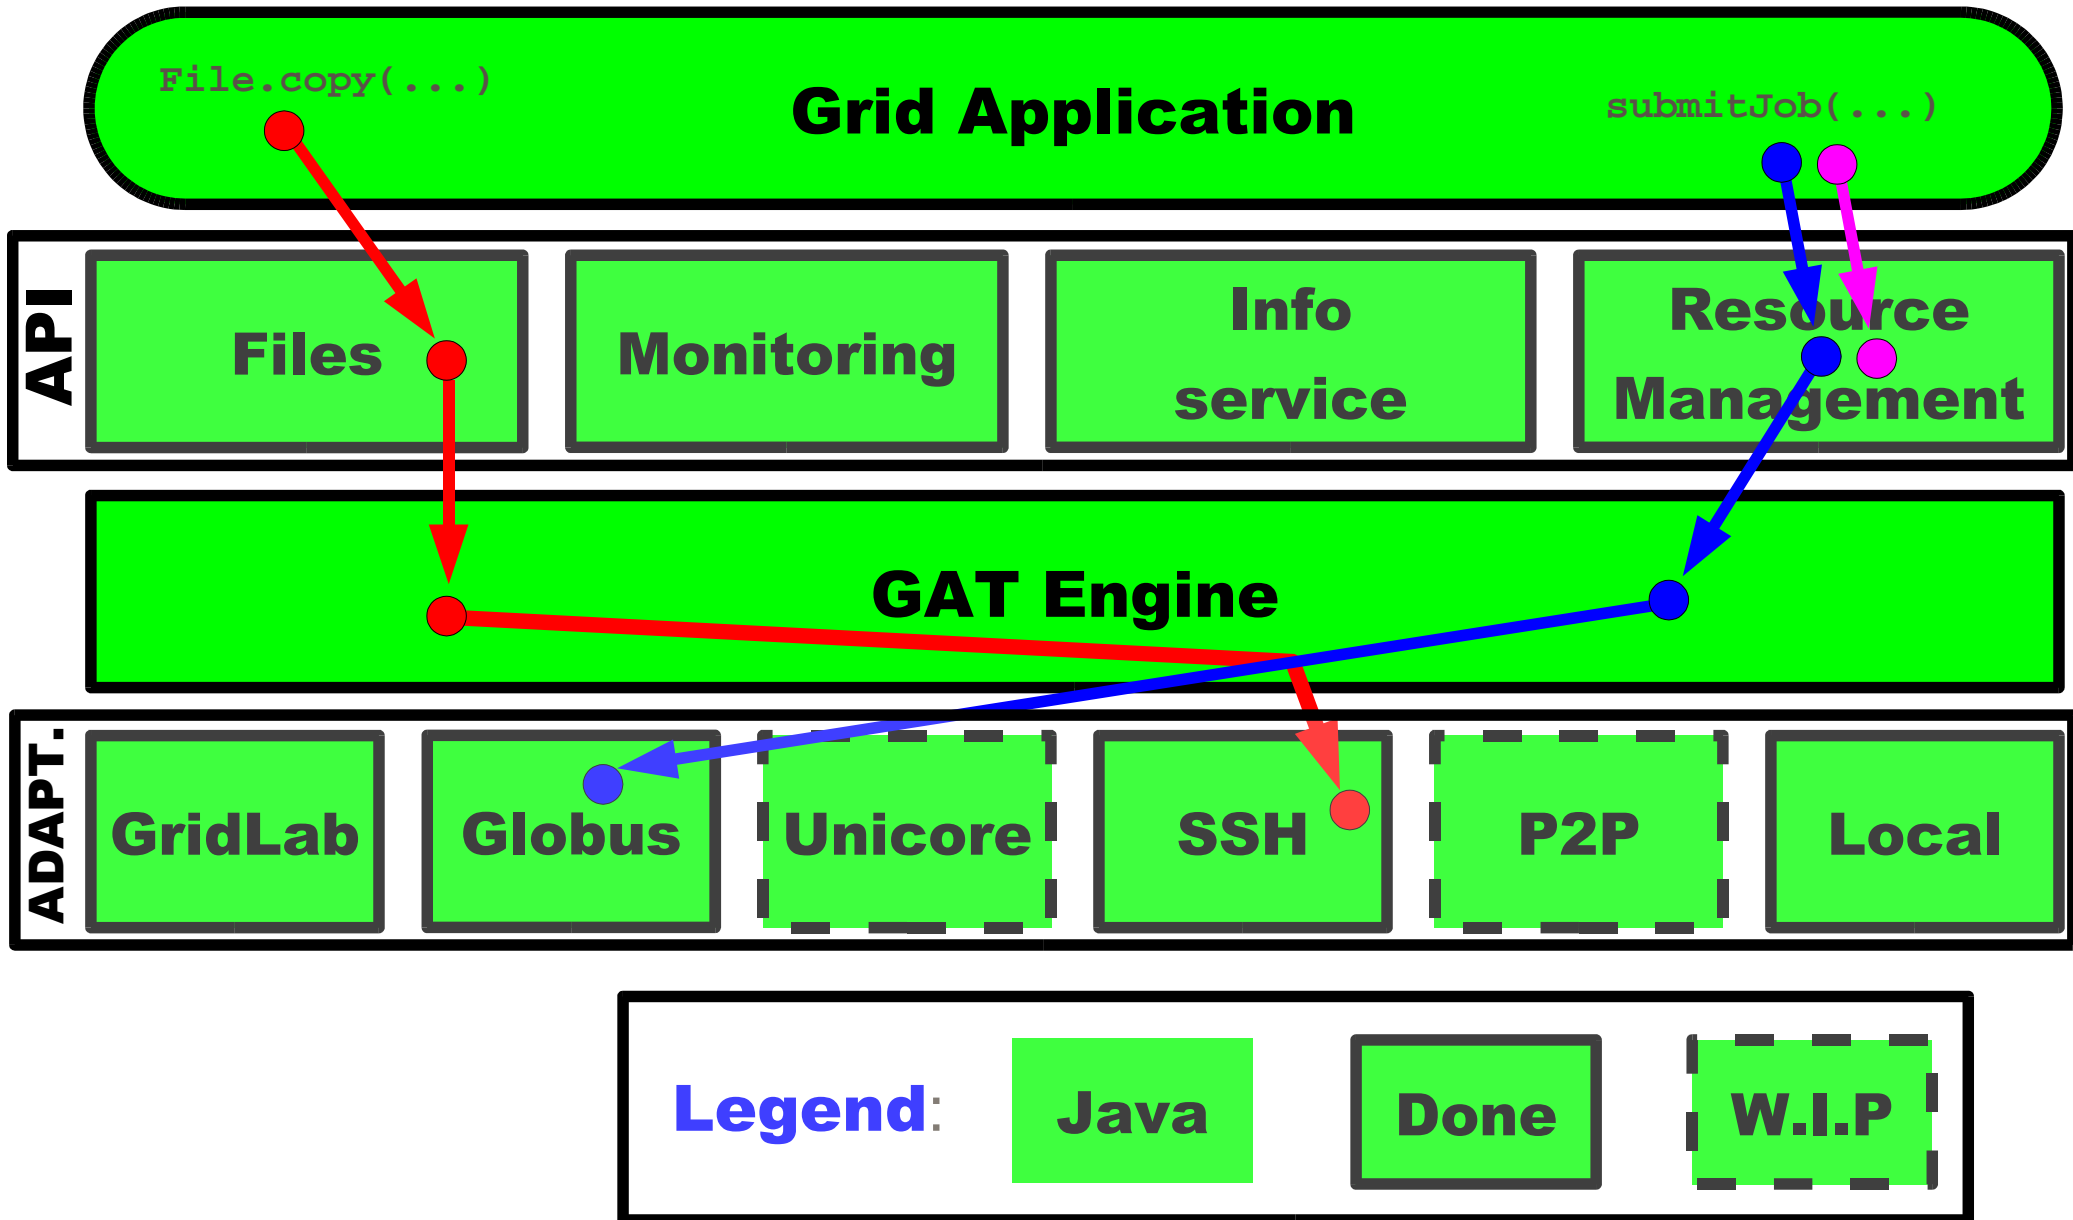


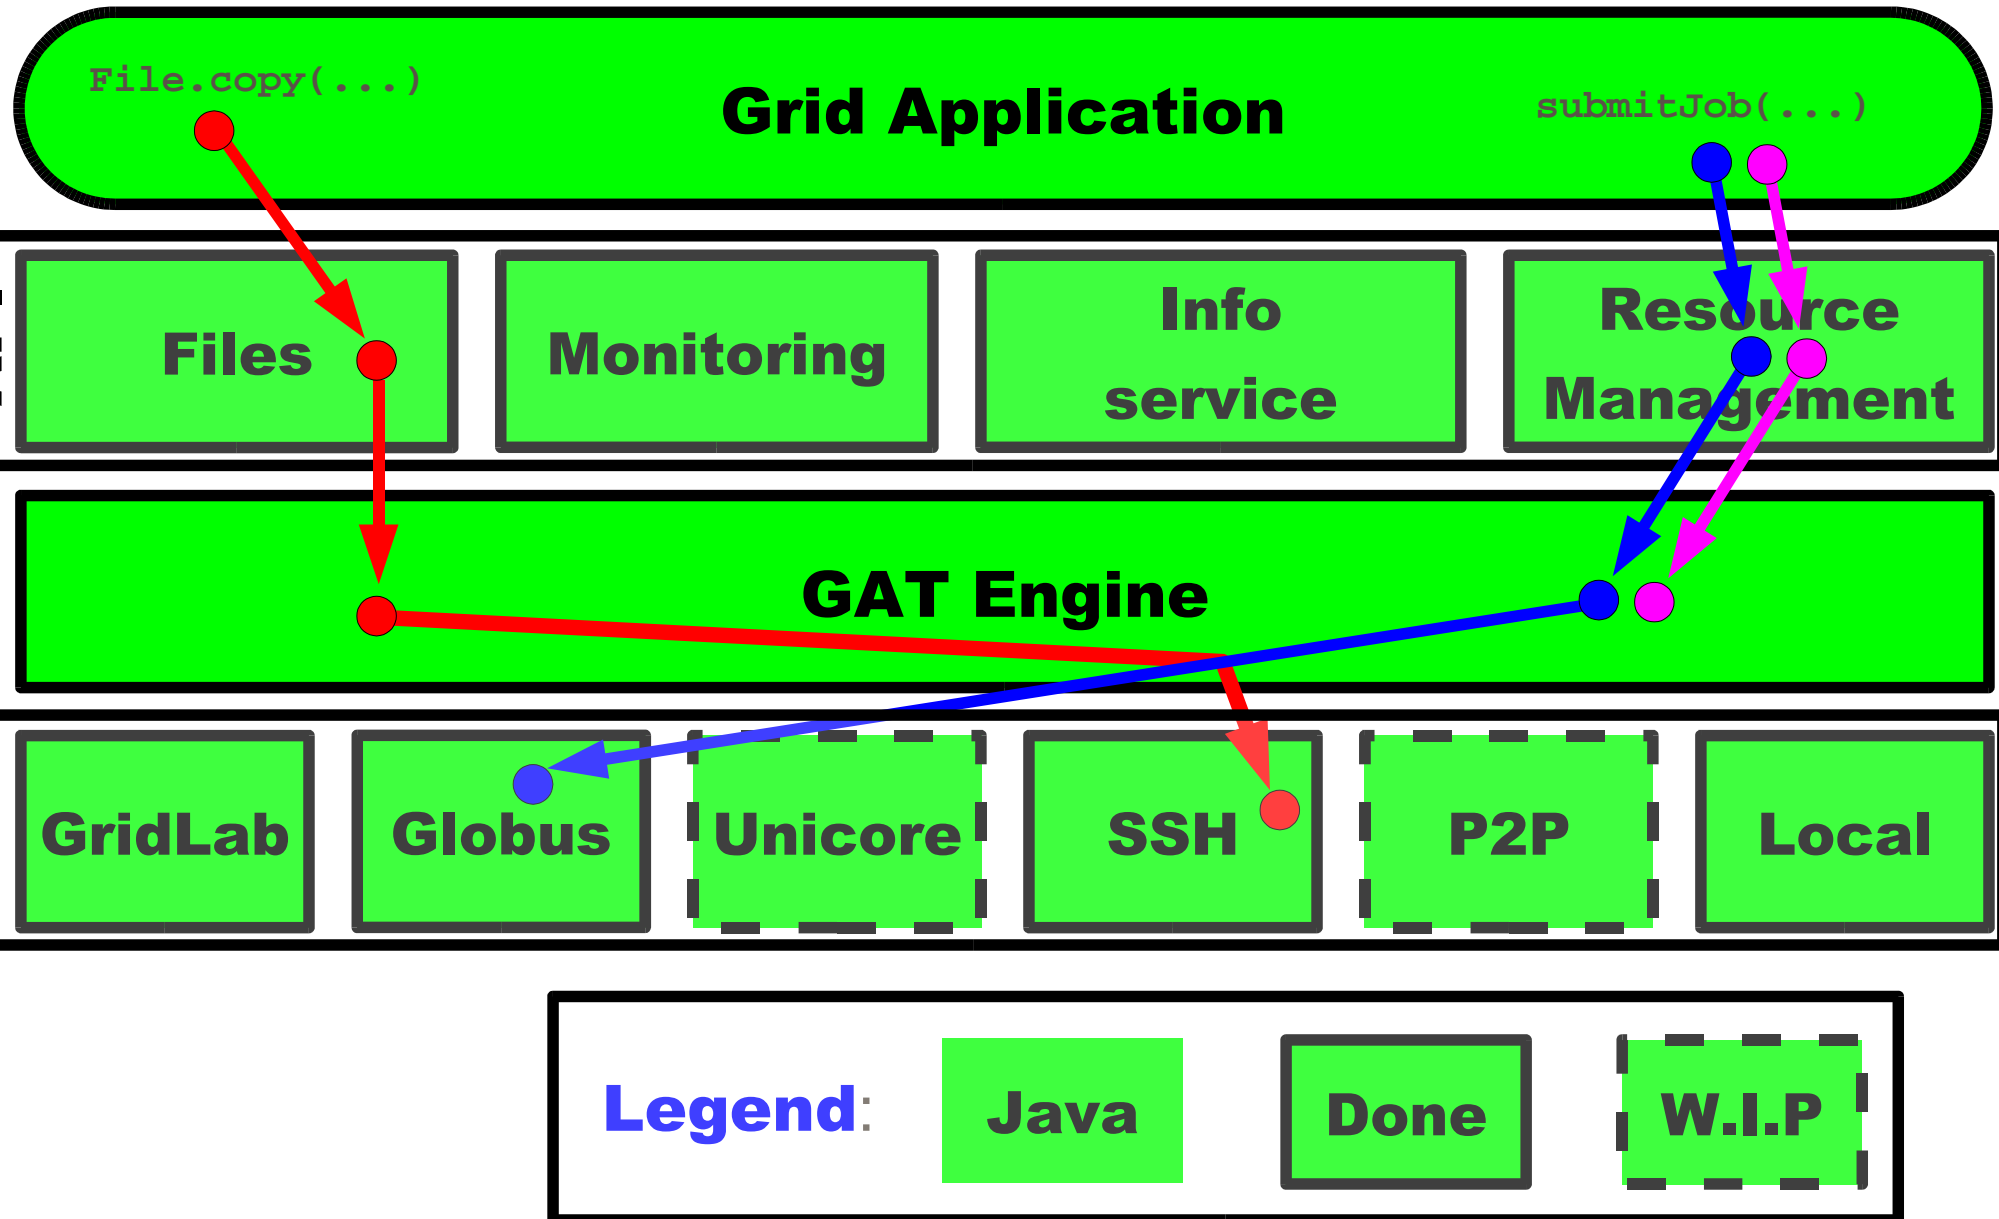


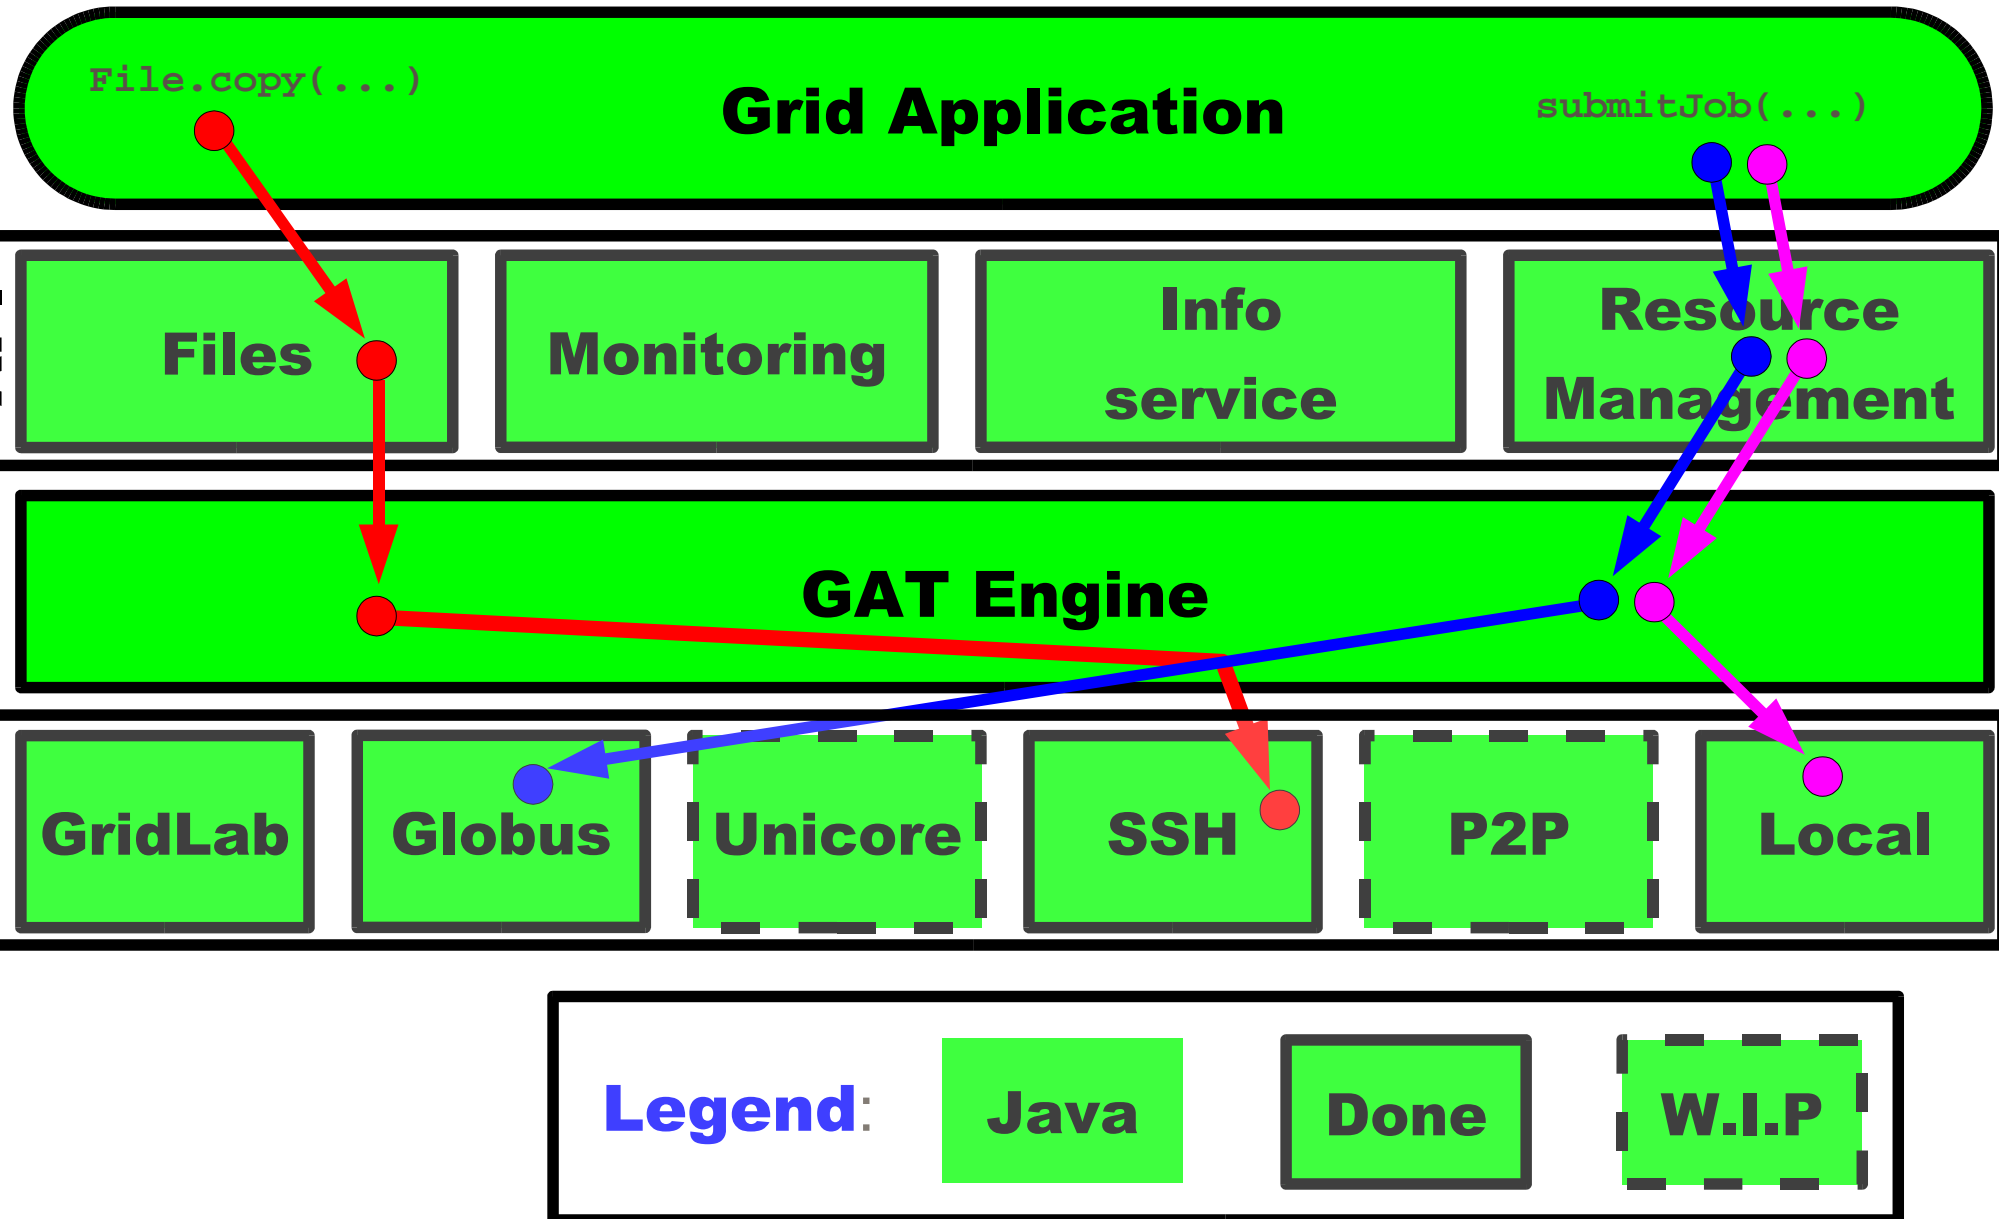












- Adaptors are Java JAR files, dynamically loaded into the application.
- Late binding (specification allows this):
 - The GAT engine selects the best adaptor for each method.
 - Example: Create file object.
 - Engine selects and instantiates all adaptors than can implement this file object.
 - File.copy from site A to site B and C.
 - A -> B copy with GridFTP.
 - A -> C copy with GridLab Data movement.
 - Provides flexibility and fault tolerance.
- Other implementations (C, C++, ...) use static binding

Spec vs JavaGAT (1)

- Memory management is done for you (GC).
- GAT utility objects are not needed (list, table).
- No status objects, but exceptions
 - Can be nested.

- File and Streaming API:
 - JavaGAT subclasses existing “java.io” classes.
 - Streams and random access are separate
 - The GAT specification defines seeks on streams.
 - Directories are “first class citizens” and can be moved, copied, etc.
 - Directories can be used with any adopter
 - Stage in and out with resource broker.
 - Any Java code that uses files can now transparently use remote files.
 - File operations (getSize, isReadable).
 - Random File access.
 - Streaming.

- Complete set of local adaptors.
- Generic:
 - Service locator (Igrid); used by GridLab adaptors
 - Stage in / out on top of GAT file; used by broker adaptors
 - Support for dealing with security contexts
 - Make writing new adaptors as easy as possible
- Many adaptors for grid middleware
 - Will be mentioned later

- What is GAT and why do we need it?
- JavaGAT structure and overview
- **Security**
- Grid I/O
- Resource Management
- Application Information Management
- Monitoring
- Hands-on session

- Preferences
 - Key value pairs
 - Adaptor-specific instructions
 - Global or local, local overrides global
 - Example: ("File.adaptor.name", "globus")
- GATContext
 - Contains security information
 - Contains global preferences
 - There can be more than one context
 - Needed to create GAT objects
- GAT
 - Factory for all GAT objects
- GAT Exceptions
 - Nested, helps debugging (Needed because of late binding)

- URI
 - Slightly different semantics compared to java.net.URI
 - Use the right number of /'s in the URI's
 - Full URI is easy....
 - `protocol://machine/<path>file`
 -but some fields may be blank
 - `file:///output` (local file in current directory)
 - `file:///output` (local file in root (/) directory)
 - `file:///tmp/output` (local file in /tmp directory)
 - `ftp://10.0.0.1/output` (remote file in default ftp dir)
- Use the right scheme (protocol) in the URI:
 - JavaGAT can choose (late binding):
 - `any://`
 - Force a specific adaptor (early binding):
 - `ftp://`, `gsiftp://`, `http://`, `file://`,

```
GATContext context = new GATContext();  
Preferences prefs = new Preferences();  
prefs.put("File.adaptor.name", "globus");  
context.addPreferences(prefs);
```

```
src = new URI("hello");  
file = GAT.createFile(context, src);
```

OR use local preferences to override globals:

```
file = GAT.createFile(context, morePrefs, src);
```

- SecurityContext
 - A container for security Information.
- Abstract, use subclasses
 - PasswordSecurityContext
 - CertificateSecurityContext
 - MyProxyServerCredentialSecurityContext
- Notes restrict the access to the context
 - Avoid broadcasting of passwords / credentials
 - Restrict access to a set of hosts or adaptors
 - One or more notes -> restricted to those adaptors/hosts
 - No notes -> any adaptor can use context for any host
- Typically not needed if default credentials / private keys are used

```
GATContext context = new GATContext();
SecurityContext pwd =
    new PasswordSecurityContext(username, password);

// restrict access
pwd.addNote("hosts", "hostname1:port1,hostname2:port2");
pwd.addNote("adaptors", "ftp,ssh");

SecurityContext cert =
    new CertificateSecurityContext(keyfile, username, passphrase);

// add them to the GAT context
context.addSecurityContext(pwd);
context.addSecurityContext(cert);
```

- What is GAT and why do we need it?
- JavaGAT structure and overview
- Security
- **Grid I/O**
- Resource Management
- Application Information Management
- Monitoring
- Hands-on session

- Copy, move, read, write files on the grid
- Random access to remote files
- Replicate files between different grid sites
- Inter-process communication

- File, streams and random file access:
 - Local files
 - GridLab data service
 - FTP, HTTP, HTTPS
 - GridFTP (Globus)
 - SSH, SFTP
- Logical file (replication):
 - Generic adaptor on top of GAT File
 - Logical file adaptor for GridLab replica service
- Pipe
 - Sockets

- File
 - FileInputStream / FileOutputStream
 - RandomAccessFile
- } Extend java.io classes
- ↓
- Grid-enable your code, just by replacing one "new" statement
- LogicalFile
 - Replicated file support
 - Basic Inter-process communication
 - Endpoint
 - Pipe
 - PileListener

- Represents both files and directories (like java.io)
 - canRead, canWrite
 - delete
 - mkdir
 - list
 - copy
 - move
 - ...

```
package tutorial;

import org.gridlab.gat.*;
import org.gridlab.gat.io.File;

public class RemoteCopy {
    public static void main(String[] args) throws Exception {
        GATContext context = new GATContext();

        URI src = new URI(args[0]);
        URI dest = new URI(args[1]);
        File file = GAT.createFile(context, src);           // create file object

        file.copy(dest);                                   // and copy it
        GAT.end();
    }
}
```




GAT File Streaming Example



```
package tutorial;  
import java.io.*;  
import org.gridlab.gat.*;  
import org.gridlab.gat.io.FileInputStream;
```

```
class RemoteCat {  
    public static void main(String[] args) throws Exception {  
        GATContext context = new GATContext();  
        URI loc = new URI(args[0]);
```

```
        FileInputStream in = GAT.createFileInputStream(context, loc);  
        InputStreamReader reader = new InputStreamReader(in);  
        BufferedReader buf = new BufferedReader(reader);
```



Standard
java.io
classes

```
        String result = buf.readLine();  
        System.err.println("result = " + result);  
        in.close();
```

```
    }
```

```
}
```

GAT Remote Random Access Files

- Random access to remote files
 - read
 - write
 - seek
 - length
 - ...

GAT Logical File (replica management)

- LogicalFile class
 - An abstract representation of a set of identical physical files
 - addFile / addURI
 - removeFile / removeURI
 - **replicate(URI destination)**
- Replicate a logical file to a new location.
 - Copy one of the files in the set to the new location
 - Choose the “best” one
 - Closest in terms of bandwidth
 - Cheapest
 - ...
 - Depends on adaptor
- Typically used for staging in files for jobs
 - Resource broker (or GAT) chooses one of the files in the set

- Basic IPC primitives
 - Not intended for parallel programming
 - API looks like socket
 - Point to point only
 - Protocol independent
- Endpoint
 - Create them
 - **listen** (bind/accept in “socket speak”)
 - **connect** (setup connection to another Endpoint)
 - **listen** and **connect** both return a Pipe
- Pipe
 - A connection between to Endpoints (like a socket)
 - Provides InputStream and OutputStream

- Endpoints are shared using the GAT application information service
- Asynchronous connection setup
- PipeListener
 - `processPipe(Pipe pipe);`
- `Endpoint.listen(PipeListener p)`
 - Asynchronous (returns immediately)
 - When a remote site connects, GAT calls `processPipe` on listener

- What is GAT and why do we need it?
- JavaGAT structure and overview
- Security
- Grid I/O
- **Resource Management**
- Application Information Management
- Monitoring
- Hands-on session

Resource Management Use Cases

- Schedule, run jobs on the grid
- Stage in / stage out files that belong with jobs
- Cancel, checkpoint, migrate jobs
- Reserve resources


Available Adaptors for Resource Management

- Local machine
- GridLab GRMS
- Globus
- ProActive
- Work-in-progress:
 - SSH
 - Resource broker adaptor for Unicore (Peggy Lindner)

- ResourceDescription
 - SoftwareResourceDescription
 - Describes a software component: library or Operating System, etc
 - Key/value pairs
 - “os.name”, “linux”
 - “os.release”, “2.6”
 - HardwareResourceDescription
 - Describes a hardware component: machine to run a job on
 - Key/value pairs
 - “memory.size”, “1GB”
 - “machine.type”, “i686”
 - Some keys and values defined in the spec, others may be adaptor specific
 - Both have a list of additional ResourceDescriptions
 - Build trees of ResourceDescriptions to express complex requirements

Two requirements

SoftwareResourceDescription
os.name=linux



HardwareResourceDescription
machine.type=i686

```
SoftwareResourceDescription srd =  
    new SoftwareResourceDescription();  
srd.addResourceAttribute("os.name", "linux");
```

```
HardwareResourceDescription hrd =  
    new HardwareResourceDescription();  
hrd.addResourceAttribute("machine.type", "i686");
```

```
srd.addResourceDescription(hrd); // require srd and hrd
```

SoftwareResourceDescription
os.name=linux

HardwareResourceDescription
machine.type=i686

HardwareResourceDescription
machine.type=power4

```
SoftwareResourceDescription srd =  
    new SoftwareResourceDescription();  
srd.addResourceAttribute("os.name", "linux");
```

```
HardwareResourceDescription hrd1 = new HardwareResourceDescription();  
hrd1.addResourceAttribute("machine.type", "i686");
```

```
HardwareResourceDescription hrd2 = new HardwareResourceDescription();  
hrd2.addResourceAttribute("machine.type", "power4");
```

```
srd.addResourceDescription(hrd1); // We must have Linux  
srd.addResourceDescription(hrd2); // on either i686 or power4
```


- Resource
 - SoftwareResource
 - HardwareResource
- Reservation
- SoftwareDescription
 - Executable
 - Arguments
 - stdin, stdout, stderr
 - Pre-staged, post-staged files
- JobDescription
 - Must have SoftwareDescription
 - Has ResourceDescription or Resource

Describes the software you want to run on the grid

Describes your job:
Software and Requirements (for brokering)
or
Software and a specific resource to run on

GAT Resource Management Classes (3)

- ResourceBroker
 - Submit job description
 - Find resources
 - Make reservations
- Job
 - Result of job submission
 - Query state
 - Cancel
 - Checkpoint
 - Migrate

Resource Management Example (1)

```
public class SubmitLocalJob {  
    public static void main(String[] args) throws Exception {  
        GATContext context = new GATContext();  
  
        SoftwareDescription sd = new SoftwareDescription();  
        sd.setLocation("file:///bin/hostname");  
        File stdout = GAT.createFile(context, "hostname.txt");  
        sd.setStdout(stdout);  
  
        JobDescription jd = new JobDescription(sd);  
        ResourceBroker broker = GAT.createResourceBroker(context);  
        Job job = broker.submitJob(jd);  
        while (job.getState() != Job.STOPPED &&  
            job.getState() != Job.SUBMISSION_ERROR) Thread.sleep(1000);  
    }  
}
```

Resource Management Example (2)

```
public class SubmitRemoteJob {
    public static void main(String[] args) throws Exception {
        // ... above code is the same as before
        ResourceDescription rd = new HardwareResourceDescription();
        rd.addAttribute("machine.node", args[0]);

        JobDescription jd = new JobDescription(sd, rd);
        ResourceBroker broker = GAT.createResourceBroker(context);
        Job job = broker.submitJob(jd);

        while (job.getState() != Job.STOPPED &&
            job.getState() != Job.SUBMISSION_ERROR) {
            System.err.println("job state = " + job.getInfo());
            Thread.sleep(1000);
        }
    }
}
```

- What is GAT and why do we need it?
- JavaGAT structure and overview
- Security
- Grid I/O
- Resource Management
- **Application Information Management**
- **Monitoring**
- **Hands-on session**

- Advert service is a directory with an hierarchical name space
- The advert service can publish GAT objects
 - File
 - LogicalFile
 - Endpoint
 - Resource (software, hardware)
 - Job
- User attaches meta data to GAT objects
- Application can query, based on meta data
- Advert service is persistent
- Advert service is language independent

- Publish your (replicated) files, so other applications can open them
- Publish Endpoints to establish inter process communication
- Publish the (hardware) resources you found
- Spawn jobs, publish their Job object, let other applications (like a portal) monitor them

Available Adaptors for the Advert Service

- Local system (file)
- GridLab Storagebox

- Advertisable
 - Marker interface (like java.io.Serializable)
 - GAT objects that implement this interface can be advertised
- MetaData
 - Table of (key,value) pairs (Strings)
 - Attached to Advertisable
- AdvertService
 - Publish adverts at some path in the hierarchical name space
 - Path uniquely identifies advert (one advert at any given path)
 - Supports current working directory (relative paths)
 - `add(advert, metaData, path)`
 - `getAdvertisable(path)`
 - `find(metaData)` metaData can contain regular expressions
 - ...

Example: Publishing GAT Objects

```
public class AdvertServicePublish {  
    public static void main(String[] args) throws Exception {  
        GATContext context = new GATContext();  
        AdvertService advert = GAT.createAdvertService(context);  
  
        File file = GAT.createFile(context, "foo");  
        Metadata meta = new Metadata();  
        meta.put("name", "myTestFile");  
        meta.put("purpose", "tutorial");  
        advert.add(file, meta, "/home/rob/tutorialFile");  
  
        Endpoint e = GAT.createEndpoint(context);  
        meta = new Metadata();  
        meta.put("name", "myTestEndpoint");  
        meta.put("purpose", "tutorial");  
        advert.add(e, meta, "/home/rob/tutorialEndpoint");  
    }  
}
```

Example: retrieving GAT Objects

```
public class AdvertServicePrinter {  
    public static void main(String[] args) throws Exception {  
        GATContext context = new GATContext();  
  
        AdvertService advert = GAT.createAdvertService(context);  
  
        Advertisable result = advert.getAdvertisable(args[0]);  
        if(result == null) {  
            System.err.println("object not found");  
            return;  
        }  
  
        System.err.println("got object back: " + result);  
    }  
}
```

- Create MetaData query description
 - All fields in query must be matched
 - Results may have additional fields
 - MetaData value fields can be regular expressions

```
/home/rob/tutorialFile  
"name", "myTestFile"  
"purpose", "tutorial"
```

```
/home/rob/tutorialEndpoint  
"name", "myTestEndpoint"  
"purpose", "tutorial"
```

- {("purpose", "tutorial")} matches both
- {("purpose", "tutorial"), ("location", "any")} matches none
- {("name", ".*File")} matches blue
- {("name", "my.*")} matches both

Example: Query the Advert Service

```
public class AdvertServiceQuery {  
    public static void main(String[] args) throws Exception {  
        GATContext context = new GATContext();  
        AdvertService advert = GAT.createAdvertService(context);  
  
        Metadata meta = new Metadata();  
        meta.put(args[0], args[1]);  
  
        String[] paths = advert.find(meta);  
        if (paths == null) { System.err.println("no objects found");return; }  
  
        System.err.println(paths.length + " adverts found");  
        for (int i = 0; i < paths.length; i++) {  
            Advertisable result = advert.getAdvertisable(paths[i]);  
            System.err.println("advert " + i + " is " + result);  
        }  
    }  
}
```

- What is GAT and why do we need it?
- JavaGAT structure and overview
- Security
- Grid I/O
- Resource Management
- Application Information Management
- **Monitoring**
- Hands-on session

GAT Monitoring and Steering System use cases

- Monitor and steer applications
 - Get application info (current iteration number)
 - Tell application to checkpoint
 - Change application variable to steer simulation
 - ...
- Monitor grid resources
 - Host load
 - Free disc space
 - Current available network bandwidth
 - ...
- Monitor and steer GAT objects
 - Monitor progress of GAT file copy
 - Monitor job status (is it in the queue or already scheduled)
 - ...

- External monitoring systems
 - Mercury
 - Monitors grid resources: hosts, queues, etc.
 - Monitors / steers applications
 - the GridLab project
 - the Datagrid project
 - the Hungarian Supergrid project
 - the P-GRADE Parallel Grid Runtime and Application Development Environment
 - Delphoi (work in progress)
 - Monitors network
 - The GridLab project
- Adaptor-specific (job management)

- Monitorable
 - A marker interface indicating which objects can be monitored
 - Interface to external monitoring system
- Monitorable GAT objects
 - File
 - LogicalFile
 - RandomAccessFile
 - FileInputStream
 - FileOutputStream
 - Endpoint
 - Pipe
 - AdvertService
 - Resource
 - Job

- MetricDefinition
 - A description of a measurable quantity within a monitoring system (free disk space)
 - Defines:
 - Measurement type: continuous / discrete
 - Data type and unit for parameters and return value
 - Create metrics by supplying parameters
 - Host name
 - Device name: /dev/hda
 - Monitoring system says:
 - I can measure free disk space.
 - I need two String parameters (device name, host name)
 - I will return a long value which is the free disk space in megabytes.
- Get list of MetricDefinitions from Monitorable

- Metric
 - Use a metric definition to define the metric you are interested in
 - MetricDefinition is: I can measure disk space
 - Metric is an **instance** of this: disk space on localhost, /dev/hda
 - Create by MetricDefinition.createMetric("localhost", "/dev/hda");
- MetricValue
 - A value measured by the monitoring system
 - Poll monitoring system with Monitorable.getMeasurement(Metric)
 - Contains time stamp, value
- MetricListener
 - Asynchronous monitoring handler, produces new MetricValues
 - Can be called whenever something changes
 - Can be called with a specific frequency (for continuous metrics)

Monitoring Example (1)

```
public class ListMetricDefinitions {  
    public static void main(String[] args) throws Exception {  
        GATContext c = new GATContext();  
        c.addPreference("monitoring.adaptor.name", "mercury");  
  
        Monitorable m = GAT.createMonitorable(c);  
  
        List definitions = m.getMetricDefinitions();  
  
        System.err.println("found " + definitions.size() + " definitions");  
  
        for (int i = 0; i < definitions.size(); i++) {  
            System.err.println(definitions.get(i));  
        }  
    }  
}
```

Monitoring Example (2)

```
public class GetMetric {  
    public static void main(String[] args) throws Exception {  
        GATContext c = new GATContext();  
        c.addPreference("monitoring.adaptor.name", "mercury");  
  
        Monitorable m = GAT.createMonitorable(c);  
        MetricDefinition def = m.getMetricDefinitionByName("host.mem.free");  
  
        HashMap params = new HashMap();  
        params.put("host", args[0]);  
  
        Metric metric = def.createMetric(params);  
        MetricValue result = m.getMeasurement(metric);  
  
        System.err.println("Free memory at " +  
            new Date(result.getEventTime()) + " was: " + result.getValue());  
    }  
}
```

The Power of GAT (1)

- Submit job to remote resource
 - Stage in a directory
 - Get output back
 - Monitor the Job
-
- Resource broker independent
 - Transfer protocol independent
-
- In 70 lines of code!


```
public class SubmitJobCallback implements MetricListener {
    boolean exit = false;

    public void start(String[] args) throws Exception {
        GATContext context = new GATContext();
        Job job = submitJob(context, args[0]);

        MetricDefinition md = job.getMetricDefinitionByName("job.status");
        Metric m = md.createMetric(); // no parameters needed for metric
        job.addMetricListener(this, m); // register callback for job.status

        // wait until job is done
        synchronized (this) {
            while (!exit) wait();
        }
    }
}

// .... continued on next slide ....
```



My `processMetricEvent` method will be called when `job.status` changes

```
public synchronized void processMetricEvent(MetricValue val) {  
    String state = (String) val.getValue();  
  
    System.err.println("SubmitJobCallback: Processing metric: "  
        + val.getMetric() + ", state is " + state);  
  
    if (state.equals("STOPPED") || state.equals("SUBMISSION_ERROR")) {  
        exit = true;  
        notifyAll();  
    }  
}  
  
// .... continued on next slide ....
```



```
Job submitJob(GATContext context, String hostname) throws Exception {
    File outFile = GAT.createFile(context, "any:///hostname.txt");
    File stageInDir = GAT.createFile(context, "any:///home/rob/testDir");

    SoftwareDescription sd = new SoftwareDescription();
    sd.setLocation("file:///bin/hostname");
    sd.setStdout(outFile);
    sd.setPreStaged(stageInDir);

    Hashtable attributes = new Hashtable();
    attributes.put("machine.node", hostname);
    ResourceDescription rd = new HardwareResourceDescription(attributes);

    JobDescription jd = new JobDescription(sd, rd);
    ResourceBroker broker = GAT.createResourceBroker(context);
    return broker.submitJob(jd);
}
```

- GAT will be supported for at least the next two years
- GAT is being standardized within GGF
 - This will take time
 - GAT 2.0 will be called SAGA (Simple API for Grid Applications)
 - SAGA is very close to GAT
 - Adds task model to GAT for asynchronous grid operations
- We will offer a migration path from GAT to SAGA

- The GAT provides a simple and stable API to various Grid environments
- But powerful!
- Independent of grid middleware
- Portable
- Downloads:
 - www.gridlab.org
 - Distributions
 - Anonymous CVS access
 - C Platforms: Linux, Windows, Mac OS X, SGI Irix, True64 UNIX
 - Java Platforms: any (Java 1.4 or higher)
 - Support via mailing list gat@gridlab.org

- Requirements:
 - Java 1.4 or newer
 - Download at java.sun.com
 - Ant 1.6 or newer
 - Java's equivalent of “make”
 - Download at ant.apache.org
- The Java GAT
 - www.gridlab.org
 - Click on Grid Application Toolkit
 - Click on GAT releases (on the left)
 - Download latest Java GAT version
 - Untar / unzip it in any place you like

- Go to the **JavaGATEngine** directory
- Type **ant**
 - This takes about 10 seconds
 - The engine and all examples are now compiled
- Go to the **JavaGATAdaptors** directory
- Type **ant**
 - This takes about 10 seconds
 - The adaptors are now compiled
- Go back to the **JavaGATEngine** directory
 - Test it:

`bin/run_gat_app tutorial.RemoteCat http://www.cs.vu.nl/~rob/hello.txt`

Or (on windows)

`bin\run_gat_app.bat tutorial.RemoteCat http://www.cs.vu.nl/~rob/hello.txt`

Capability Provider Interface (CPI)

- GAT objects are Java interfaces
- Below the API level, the GAT Engine provides CPI classes (Capability Provider Interface)
 - CPI classes implement the GAT object interfaces
 - All methods in the CPI classes throw exceptions
 - CPI classes provide additional generic code
- Adaptors extend a CPI class
 - Fill in some or all methods
 - Late binding: if a method is not implemented GAT selects another adaptor

